

Orthopedic Patients Classification

Chhaya Mudgal

12/7/2020

Introduction and Project Outline

This script is the Capstone Own Project for HarvardX Data Science Professional Certificate. Dataset chosen to study classification of Orthopedic Patients based on Biomechanical Features. Data has been downloaded from Kaggle database. In this project KNN (Nearest Neighbour Algorithm) Supported Vector Machine (SVM) and Random Forest will be used for comparing accuracy of patients classification..

Machine Learning is being used in various medical fields to predict and classify diseases. Orthopedic health condition of a patient can be detected from the biomechanical features. Application of machine learning algorithms in medical science helps in classification. Different algorithms are applied to detect diseases and classify patients accordingly. In this project various machine learning algorithms are applied to find out which one works most accurately to detect and classify orthopedic patients. Algorithms compared for accuracy are KNN, SVM and Random Forest. Each of the patients in the dataset is represented by six biomechanical attributes derived from the shape and orientation of pelvis and lumbar spine.

Why have I chose the data? I have chosen this dataset because: It is freely available online on Kaggle. It is 'medium' sized. Not small or too big to be processed on my personal computer. There are a reasonable number of predictor columns, and easy to understand.

Load Data

Step1: Load Library Packages

```
#Step1: Load Library Packages
if(!require(tidyverse)) install.packages("tidyverse", repos = "http://cran.us.r-project.org")
if(!require(caret)) install.packages("caret", repos = "http://cran.us.r-project.org")
if(!require(data.table)) install.packages("data.table", repos = "http://cran.us.r-project.org")
if(!require(corrplot)) install.packages("corrplot", repos = "http://cran.us.r-project.org")
if(!require(ggplot2)) install.packages("ggplot2", repos = "http://cran.us.r-project.org")
if(!require(gmodels)) install.packages("gmodels", repos = "http://cran.us.r-project.org")
if(!require(knn)) install.packages("knn", repos = "http://cran.us.r-project.org")

library(gmodels)
library(tidyverse)
library(caret)
library(data.table)
library(corrplot)
library(ggplot2)
library(class)
```

Step2: Load Data

Data is downloaded from file stored in github repository Initial Step is to know the data set, hence the peek at the top 10 and last 10 lines for the dataset. Data set has 310 rows and 7 columns. Analyzing the structure of the dataset it is found that the six features have numeric values and the 7th column is the label that tells whether the patient falls under normal or abnormal category. This label is of type factor.

```
# Step2: Load Data
# Data file is stored in github for easy
# Getting data file from https://github.com/cmudgal/DataScienceHarvardX/tree/master/CapstoneProject/OrthoPatientData

u="https://raw.githubusercontent.com/cmudgal/DataScienceHarvardX/master/CapstoneProject/OrthoPatientData/ortho_data.csv"
ortho_data <- read.csv(u) # "Column_2C_weka.csv"

# Peek at first 10 rows of data
head(ortho_data, n=10)
```

```
##      pelvic_incidence pelvic_tilt.numeric lumbar_lordosis_angle sacral_slope
## 1          63.02782         22.552586          39.60912         40.47523
## 2          39.05695         10.060991          25.01538         28.99596
## 3          68.83202         22.218482          50.09219         46.61354
## 4          69.29701         24.652878          44.31124         44.64413
## 5          49.71286          9.652075          28.31741         40.06078
## 6          40.25020         13.921907          25.12495         26.32829
## 7          53.43293         15.864336          37.16593         37.56859
## 8          45.36675         10.755611          29.03835         34.61114
## 9          43.79019         13.533753          42.69081         30.25644
## 10         36.68635          5.010884          41.94875         31.67547
##      pelvic_radius degree_spondylolisthesis      class
## 1          98.67292          -0.2544000 Abnormal
## 2         114.40543           4.5642586 Abnormal
## 3         105.98514          -3.5303173 Abnormal
## 4         101.86850          11.2115234 Abnormal
## 5         108.16872           7.9185006 Abnormal
## 6         130.32787           2.2306517 Abnormal
## 7         120.56752           5.9885507 Abnormal
## 8         117.27007          -10.6758708 Abnormal
## 9         125.00289          13.2890182 Abnormal
## 10         84.24142           0.6644371 Abnormal
```

```
# Peek at last 10 rows of data
tail(ortho_data, n=10)
```

```
##      pelvic_incidence pelvic_tilt.numeric lumbar_lordosis_angle sacral_slope
## 301          50.67668           6.461501          35.00000         44.21518
## 302          89.01488          26.075981          69.02126         62.93889
## 303          54.60032          21.488974          29.36022         33.11134
## 304          34.38230           2.062683          32.39082         32.31962
## 305          45.07545          12.306951          44.58318         32.76850
## 306          47.90357          13.616688          36.00000         34.28688
## 307          53.93675          20.721496          29.22053         33.21525
```

```
## 308      61.44660      22.694968      46.17035      38.75163
## 309      45.25279      8.693157      41.58313      36.55963
## 310      33.84164      5.073991      36.64123      28.76765
##      pelvic_radius degree_spondylolisthesis class
## 301      116.5880      -0.2147106 Normal
## 302      111.4811      6.0615084 Normal
## 303      118.3433      -1.4710673 Normal
## 304      128.3002      -3.3655156 Normal
## 305      147.8946      -8.9417094 Normal
## 306      117.4491      -4.2453954 Normal
## 307      114.3658      -0.4210104 Normal
## 308      125.6707      -2.7078795 Normal
## 309      118.5458      0.2147502 Normal
## 310      123.9452      -0.1992491 Normal
```

```
# Get number of rows and columns in data set
# Data set has 310 rows and 7 columns
dim(ortho_data)
```

```
## [1] 310 7
```

```
# Get summary of the data
summary(ortho_data)
```

```
## pelvic_incidence pelvic_tilt.numeric lumbar_lordosis_angle sacral_slope
## Min. : 26.15 Min. : -6.555 Min. : 14.00 Min. : 13.37
## 1st Qu.: 46.43 1st Qu.: 10.667 1st Qu.: 37.00 1st Qu.: 33.35
## Median : 58.69 Median : 16.358 Median : 49.56 Median : 42.40
## Mean : 60.50 Mean : 17.543 Mean : 51.93 Mean : 42.95
## 3rd Qu.: 72.88 3rd Qu.: 22.120 3rd Qu.: 63.00 3rd Qu.: 52.70
## Max. : 129.83 Max. : 49.432 Max. : 125.74 Max. : 121.43
## pelvic_radius degree_spondylolisthesis class
## Min. : 70.08 Min. : -11.058 Abnormal:210
## 1st Qu.: 110.71 1st Qu.: 1.604 Normal :100
## Median : 118.27 Median : 11.768
## Mean : 117.92 Mean : 26.297
## 3rd Qu.: 125.47 3rd Qu.: 41.287
## Max. : 163.07 Max. : 418.543
```

```
# Gives the structure of data info
# All non null data set with 7 variables
# 6 are numeric features
# class is the factor with levels "Abnormal" and "Normal"
str(ortho_data)
```


```
## 'data.frame': 310 obs. of 7 variables:
## $ pelvic_incidence : num 63 39.1 68.8 69.3 49.7 ...
## $ pelvic_tilt.numeric : num 22.55 10.06 22.22 24.65 9.65 ...
## $ lumbar_lordosis_angle : num 39.6 25 50.1 44.3 28.3 ...
## $ sacral_slope : num 40.5 29 46.6 44.6 40.1 ...
## $ pelvic_radius : num 98.7 114.4 106 101.9 108.2 ...
## $ degree_spondylolisthesis: num -0.254 4.564 -3.53 11.212 7.919 ...
## $ class : Factor w/ 2 levels "Abnormal","Normal": 1 1 1 1 1 1 1 1 1 ...
```

Explore Data

There are highly positive correlations between Pelvic Incidence and Sacral Slope , also, between Pelvic Incidence and Lumbar Lordosis Angle as can be seen by scatter plots below. From the figure, as it seems Normal class values are smaller than Abnormal values; therefore narrowed with selecting some correlated features. Lets look at correlation matrix of the select features: pelvic_raduis, pelvic_incidence and lumbar_lordosis_angle.

There are 210 patients in Abnormal and 100 in Normal category

```
# Explore Data for Orthopedic Patients
plot(ortho_data$class, freq=TRUE, col="red", border="white",
     main="Distribution of patients", xlab="Class", ylab="Count")
```



```
#
# Divison of 'class' attribute of the patients
table(ortho_data$class)
```


```
##
## Abnormal   Normal
##      210      100
```

```
# Percentual division of patients using the 'class' attribute
round(prop.table(table(ortho_data$class)) * 100, digits = 1)
```


```
##
## Abnormal    Normal
##      67.7      32.3
```

```
# Correlation plot
```

```
com = ortho_data[,1:6]
cc = cor(com, method = "spearman")
corrplot(cc, tl.col = "black", order = "hclust", hclust.method = "average", addrect = 4, tl.cex = 0.7)
```




```
# Scatter plot for pelvic_radius and sacral_slope for distribution of patients
ggplot(ortho_data, aes(x = pelvic_radius, y = sacral_slope)) +
  geom_point(aes(color = factor(class)))
```




Scatter plot for pelvic_incidence and sacral_slope for distribution of patients

```
ggplot(ortho_data, aes(x = pelvic_incidence, y = sacral_slope)) +  
  geom_point(aes(color = factor(class)))
```

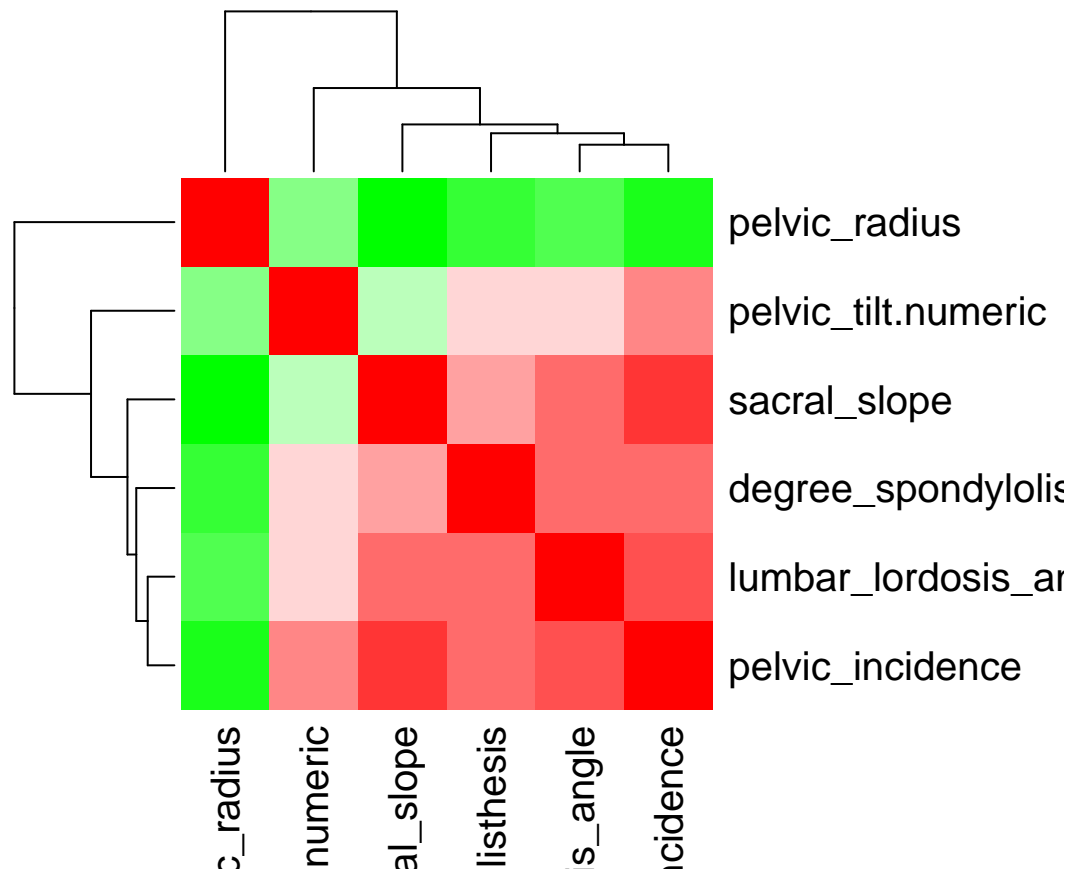


Scatter plot for pelvic_incidence and Lumbar Lordosis Angle for distribution of patients

```
ggplot(ortho_data, aes(x = pelvic_incidence, y = lumbar_lordosis_angle)) +  
  geom_point(aes(color = factor(class)))
```



```
#heat map  
palette = colorRampPalette(c("green", "white", "red")) (20)  
heatmap(x = cc, col = palette, symm = TRUE)
```

Prepare Data

This step involves, cleaning, normalizing and splicing of data. The Biomechanical Orthopedic data set will be used for classification, which is an example of predictive modeling. The last attribute of the data set, class, will be the target variable or the variable that I want to predict. 1) Check of null values 2) normalize data 3) Split data in test and train data sets

Normalize data: Looking at the summary output it is seen that all the features are not in consistent range. Look at the minimum and maximum values of all the (numerical) attributes. If one attribute has a wide range of values, need to normalize the dataset, because this means that the distance will be dominated by this feature. In the current dataset it is degree_spondylolisthesis that has wide range from -11.058 to 418.543

In order to assess the performance of the model data set is divided into two parts: a training set and a test set. The first is used to train the system, while the second is used to evaluate the learned or trained system. 70% of the original data set is as the training set, while the 10% that remains will compose the test set.

```
# check if isna
data_na<-apply(ortho_data, 2, function(x) any(is.na(x)))
data_na
```

```
##      pelvic_incidence      pelvic_tilt.numeric      lumbar_lordosis_angle
##              FALSE              FALSE              FALSE
##      sacral_slope      pelvic_radius      degree_spondylolisthesis
##              FALSE              FALSE              FALSE
```

```
##          class
##          FALSE
```

```
# There is no na in the dataset
```

```
# Build normalize function
```

```
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x))) }
```

```
# Apply normalize function to the orthopedic data set
```

```
ortho_data.norm<-as.data.frame(lapply(ortho_data[,c(1,2,3,4,5,6)], normalize))
```

```
# Summary for normalized data
```

```
summary(ortho_data.norm)
```

```
## pelvic_incidence pelvic_tilt.numeric lumbar_lordosis_angle sacral_slope
## Min. :0.0000 Min. :0.0000 Min. :0.0000 Min. :0.0000
## 1st Qu.:0.1956 1st Qu.:0.3076 1st Qu.:0.2058 1st Qu.:0.1849
## Median :0.3139 Median :0.4093 Median :0.3183 Median :0.2687
## Mean :0.3313 Mean :0.4304 Mean :0.3394 Mean :0.2738
## 3rd Qu.:0.4507 3rd Qu.:0.5122 3rd Qu.:0.4385 3rd Qu.:0.3639
## Max. :1.0000 Max. :1.0000 Max. :1.0000 Max. :1.0000
## pelvic_radius degree_spondylolisthesis
## Min. :0.0000 Min. :0.00000
## 1st Qu.:0.4369 1st Qu.:0.02947
## Median :0.5182 Median :0.05313
## Mean :0.5145 Mean :0.08695
## 3rd Qu.:0.5956 3rd Qu.:0.12185
## Max. :1.0000 Max. :1.00000
```

```
# Split Data into Training and Test Sets
```

```
# To make training and test sets, set a seed. This is a number of R's random number generator.
```

```
# The major advantage of setting a seed is that it gives same sequence of random numbers.
```

```
set.seed(123)
```

```
ortho_data.ind <- sample(1:nrow(ortho_data.norm),size=nrow(ortho_data.norm)*0.7,replace = FALSE) #random
```

```
ortho_data.train <- ortho_data.norm[ortho_data.ind,] # 70% training data
```

```
# Inspect training set
```

```
head(ortho_data.train)
```

```
## pelvic_incidence pelvic_tilt.numeric lumbar_lordosis_angle sacral_slope
## 179 0.5256865 0.5876264 0.4196986 0.3788822
## 14 0.2644930 0.4825382 0.1709289 0.1827127
## 195 0.4484296 0.4551798 0.4832544 0.3733742
## 306 0.2098221 0.3602926 0.1968814 0.1935909
## 118 0.5776410 0.8092195 0.3031205 0.3139261
## 299 0.3892446 0.4903408 0.1586459 0.2983693
## pelvic_radius degree_spondylolisthesis
## 179 0.5379260 0.14787138
## 14 0.4396688 0.04213903
```

```
## 195      0.5041576      0.08482849
## 306      0.5093802      0.01585839
## 118      0.5593208      0.17003315
## 299      0.6325551      0.02927222
```

```
ortho_data.test <- ortho_data.norm[-ortho_data.ind,] # remaining 30% test data

# Inspect test set
head(ortho_data.test)
```

```
##      pelvic_incidence pelvic_tilt.numeric lumbar_lordosis_angle sacral_slope
## 2      0.12450104      0.2967831      0.09857833  0.144629352
## 3      0.41166648      0.5139323      0.32299466  0.307660537
## 8      0.18535588      0.3091900      0.13458053  0.196591648
## 12     0.04903709      0.4335087      0.01342373  0.001384728
## 15     0.30044817      0.5491263      0.29532214  0.182712658
## 18     0.04945783      0.1732483      0.16612314  0.136628277
##      pelvic_radius degree_spondylolisthesis
## 2      0.4766489      0.0363649708
## 3      0.3860969      0.0175229033
## 8      0.5074553      0.0008899132
## 12     0.5374088      0.0269038552
## 15     0.5024710      0.0391645164
## 18     0.6337221      0.0341740122
```

```
# Compose 'class' training labels
ortho_data.trainLabels <- ortho_data[ortho_data.ind,7]

# Inspect result
#print(ortho_data.trainLabels)

# Compose 'class' test labels
ortho_data.testLabels <- ortho_data[-ortho_data.ind, 7]

# Inspect result
#print(ortho_data.testLabels)
```

Models/ Algorithms and Evaluation

KNN K Nearest Neighbour Algorithm

Build Classifier to find the k nearest neighbour for the training set using the knn() function, which uses the Euclidian distance measure to find the k-nearest neighbours to the new instance. KNN model is done in 2 ways using caret and class package. In Class package, We have to decide on the number of neighbors (k). There are several rules of thumb, one being the square root of the number of observations in the training set. In this case, we select 16 as the number of neighbors, which is approximately the square root of our sample size $N = 217$. Infact the model was run for both $k=16$ and $K=17$.

In caret package, the function picks the optimal number of neighbors (k) for you.

```
# To find the k parameter for the knn function
nr<-NROW(ortho_data.trainLabels)
```

```
# sqrt of 217 is 14.7
sqrt(nr)
```

```
## [1] 14.73092
```

```
dim(ortho_data.trainLabels)
```

```
## NULL
```

```
dim(ortho_data.train)
```

```
## [1] 217 6
```

```
ortho_data.knn.15 <- knn(train=ortho_data.train, test=ortho_data.test, cl=ortho_data.trainLabels, k=14)
ortho_data.knn.16 <- knn(train=ortho_data.train, test=ortho_data.test, cl=ortho_data.trainLabels, k=15)
```

```
# Inspect
ortho_data.knn.15
```

```
## [1] Abnormal Abnormal Abnormal Abnormal Abnormal Normal Abnormal Abnormal
## [9] Abnormal Normal Abnormal Abnormal Normal Abnormal Abnormal Abnormal
## [17] Abnormal Abnormal Abnormal Abnormal Abnormal Abnormal Normal Abnormal
## [25] Abnormal Abnormal Abnormal Abnormal Abnormal Normal Abnormal Abnormal
## [33] Abnormal Normal Abnormal Abnormal Abnormal Abnormal Abnormal Abnormal
## [41] Abnormal Abnormal Abnormal Abnormal Abnormal Abnormal Abnormal Abnormal
## [49] Normal Abnormal Abnormal Abnormal Abnormal Abnormal Abnormal Abnormal
## [57] Abnormal Abnormal Abnormal Abnormal Normal Abnormal Normal Normal
## [65] Abnormal Normal Abnormal Normal Normal Normal Normal Abnormal
## [73] Normal Abnormal Normal Normal Normal Normal Normal Abnormal
## [81] Normal Normal Normal Normal Normal Abnormal Normal Abnormal
## [89] Normal Abnormal Abnormal Abnormal Normal
## Levels: Abnormal Normal
```

```
# ortho_data.knn.15 stores the knn() function that takes as arguments
# the training set, the test set, the train labels and the amount of
# neighbours to find with this algorithm. The result of this function
# is a factor vector with the predicted classes for each row of the test data.
```

```
# Note that the test labels will be used to see if the model is good at prediction.
```

```
# Model Evaluation
```

```
# An essential next step in machine learning is the evaluation
# of the model's performance. Analyze the degree of correctness of the model's predictions.
```

```
# Put 'ortho_data.testLabels' in a data frame
```

```
ortho_dataTestLabels <- data.frame(ortho_data.testLabels)
```

```
# Merge 'ortho_data.knn.15' and 'ortho_data.testLabels'
```

```
merge <- data.frame(ortho_data.knn.16, ortho_data.testLabels)
```

```
# Inspect 'merge'
#merge
```

```
CrossTable(x = ortho_data.testLabels, y = ortho_data.knn.16, prop.chisq=FALSE)
```

```
##
##
##      Cell Contents
## |-----|
## |              N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table:  93
##
##
##               | ortho_data.knn.16
## ortho_data.testLabels | Abnormal | Normal | Row Total |
## -----|-----|-----|-----|
##           Abnormal |      54 |      8 |      62 |
##                   |  0.871 |  0.129 |  0.667 |
##                   |  0.857 |  0.267 |
##                   |  0.581 |  0.086 |
## -----|-----|-----|
##           Normal |      9 |     22 |      31 |
##                   |  0.290 |  0.710 |  0.333 |
##                   |  0.143 |  0.733 |
##                   |  0.097 |  0.237 |
## -----|-----|-----|
##           Column Total |      63 |      30 |      93 |
##                   |  0.677 |  0.323 |
## -----|-----|-----|
##
##
```

```
##create confusion matrix
```

```
tab <- table(ortho_data.knn.16, ortho_data.testLabels)
```

```
##this function divides the correct predictions by total number of predictions that tell us how accurat
```

```
accuracy <- function(x){sum(diag(x)/(sum(rowSums(x)))) * 100}
accuracy(tab)
```

```
## [1] 81.72043
```

```
#Calculate the proportion of correct classification for k = 15,16
ACC.15 <- 100 * sum( ortho_data.testLabels == ortho_data.knn.15)/NROW( ortho_data.testLabels)
ACC.16 <- 100 * sum( ortho_data.testLabels == ortho_data.knn.16)/NROW( ortho_data.testLabels)
ACC.15
```

```
## [1] 80.64516
```

```
ACC.16
```

```
## [1] 81.72043
```

```
# confusion Matrix
confusionMatrix(table(ortho_data.knn.16, ortho_data.testLabels))
```

```
## Confusion Matrix and Statistics
##
##               ortho_data.testLabels
## ortho_data.knn.16 Abnormal Normal
##           Abnormal      54      9
##           Normal       8     22
##
##           Accuracy : 0.8172
##           95% CI : (0.7235, 0.8898)
##           No Information Rate : 0.6667
##           P-Value [Acc > NIR] : 0.0009547
##
##           Kappa : 0.5854
##
##  Mcnemar's Test P-Value : 1.0000000
##
##           Sensitivity : 0.8710
##           Specificity : 0.7097
##           Pos Pred Value : 0.8571
##           Neg Pred Value : 0.7333
##           Prevalence : 0.6667
##           Detection Rate : 0.5806
##           Detection Prevalence : 0.6774
##           Balanced Accuracy : 0.7903
##
##           'Positive' Class : Abnormal
##
```

```
##### Using Caret Package
# Create index to split based on labels
index <- createDataPartition(ortho_data$class, p=0.7, list=FALSE)
# Subset training set with index
ortho.training <- ortho_data[index,]

# Subset test set with index
ortho.test <- ortho_data[-index,]

# Overview of algos supported by caret
```

```

#names(getModelInfo())

# Train a model
model_knn <- train(ortho.training[, 1:6], ortho.training[, 7], method='knn')

# Predict the labels of the test set
predictions_knn<-predict(object=model_knn,ortho.test[,1:6])

# Evaluate the predictions
table(predictions_knn)

```

```

## predictions_knn
## Abnormal    Normal
##          63         30

```

```

# Confusion matrix kNN
confusionMatrix(predictions_knn,ortho.test[,7])

```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Abnormal Normal
##   Abnormal      54      9
##   Normal        9     21
##
##              Accuracy : 0.8065
##              95% CI : (0.7115, 0.8811)
##   No Information Rate : 0.6774
##   P-Value [Acc > NIR] : 0.004057
##
##              Kappa : 0.5571
##
##  Mcnemar's Test P-Value : 1.000000
##
##              Sensitivity : 0.8571
##              Specificity : 0.7000
##              Pos Pred Value : 0.8571
##              Neg Pred Value : 0.7000
##              Prevalence : 0.6774
##              Detection Rate : 0.5806
##   Detection Prevalence : 0.6774
##   Balanced Accuracy : 0.7786
##
##              'Positive' Class : Abnormal
##

```

SVM Linear (Support Vector Machine)

support vector machines are supervised learning models with associated learning algorithms that analyze data used for classification and regression analysis. They are mostly used in classification problems.

```

# Train a model
model_svm <- train(ortho.training[, 1:6], ortho.training[, 7],method='svmLinear',trControl=trainControl

# Predict the labels of the test set
predictions_svm<-predict(object=model_svm,ortho.test[,1:6])

# Evaluate the predictions
table(predictions_svm)

## predictions_svm
## Abnormal    Normal
##          65         28

# Confusion matrix SVM
confusionMatrix(predictions_svm,ortho.test[,7])

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Abnormal Normal
##   Abnormal      56      9
##   Normal        7     21
##
##              Accuracy : 0.828
##              95% CI : (0.7357, 0.8983)
##   No Information Rate : 0.6774
##   P-Value [Acc > NIR] : 0.0008401
##
##              Kappa : 0.5994
##
##  Mcnemar's Test P-Value : 0.8025873
##
##              Sensitivity : 0.8889
##              Specificity : 0.7000
##              Pos Pred Value : 0.8615
##              Neg Pred Value : 0.7500
##              Prevalence : 0.6774
##              Detection Rate : 0.6022
##   Detection Prevalence : 0.6989
##   Balanced Accuracy : 0.7944
##
##              'Positive' Class : Abnormal
##

```

Random Forest Algorithm

Random Forest is one such very powerful ensembling machine learning algorithm. It works by creating multiple decision trees and combining the output generated by each of the decision trees. Decision tree is a classification model which works on the concept of information gain at every node.


```

# Train a model
model_rf <- train(ortho.training[, 1:6], ortho.training[, 7], method='rf')

# Predict the labels of the test set
predictions_rf<-predict(object=model_rf,ortho.test[,1:6])

# Evaluate the predictions
table(predictions_rf)

```

```

## predictions_rf
## Abnormal   Normal
##         62      31

```

```

# Confusion matrix Random Forest
confusionMatrix(predictions_rf,ortho.test[,7])

```

```

## Confusion Matrix and Statistics
##
##              Reference
## Prediction Abnormal Normal
##   Abnormal      53      9
##   Normal       10     21
##
##              Accuracy : 0.7957
##              95% CI : (0.6995, 0.8723)
##   No Information Rate : 0.6774
##   P-Value [Acc > NIR] : 0.008095
##
##              Kappa : 0.5366
##
##  Mcnemar's Test P-Value : 1.000000
##
##              Sensitivity : 0.8413
##              Specificity : 0.7000
##              Pos Pred Value : 0.8548
##              Neg Pred Value : 0.6774
##              Prevalence : 0.6774
##              Detection Rate : 0.5699
##   Detection Prevalence : 0.6667
##   Balanced Accuracy : 0.7706
##
##              'Positive' Class : Abnormal
##

```


Analysis and Conclusion

After comparing KNN, Supported Vector Machine (SVM) and Random Forest Algorithms the prediction for this data set is highest for KNN algorithm. The results of the Cross Table indicate that our model did not predict mother's job very well. To read the Cross Table, we begin by examining the top-left to bottom-right diagonal of the matrix. The diagonal of the matrix represents the number of cases that were correctly classified for each category. If the model correctly classified all cases, the matrix would have zeros

everywhere but the diagonal. In this case, we see that the numbers are quite high in the off-diagonals, indicating that our model did not successfully classify our outcome based on our predictors.

Confusion matrix or error matrix is used for summarizing the performance of a classification algorithm. Calculating a confusion matrix gives an idea of where the classification model is right and what types of errors it is making. A confusion matrix is used to check the performance of a classification model on a set of test data for which the true values are known. It can be seen that random forest performed with 79%, knn with 80% accuracy and svm linear with 82% accuracy.

```
method <- c('kNN', 'svm', 'random forest')
accuracy <- c(0.8065, 0.828, 0.79577)
df <- data.frame(method, accuracy)
ggplot(data=df, aes(x=method, y=accuracy)) +
  geom_bar(stat="identity", fill="steelblue") +
  geom_text(aes(label=accuracy), vjust=-0.3, size=3.5) +
  theme_minimal()
```



Next Steps.. Optimization:

For kNN algorithm, the tuning parameters are 'k' value and number of 'features/attributes selection. Optimum 'k' value can be found using graph below. It was found that max accuracy in knn is at k=21. It increases to about 82%

```
# Optimization
i=1
```

```

k.optm=1
for (i in 1:28){
  knn.mod <- knn(train=ortho_data.train, test=ortho_data.test, cl=ortho_data.trainLabels, k=i)
  k.optm[i] <- 100 * sum( ortho_data.testLabels == knn.mod)/NROW( ortho_data.testLabels)
  k=i
  cat(k, '=', k.optm[i], '
',)
}

```

```

## 1 = 78.49462
## 2 = 73.11828
## 3 = 74.19355
## 4 = 74.19355
## 5 = 74.19355
## 6 = 77.41935
## 7 = 77.41935
## 8 = 72.04301
## 9 = 78.49462
## 10 = 79.56989
## 11 = 80.64516
## 12 = 79.56989
## 13 = 79.56989
## 14 = 79.56989
## 15 = 81.72043
## 16 = 79.56989
## 17 = 80.64516
## 18 = 81.72043
## 19 = 81.72043
## 20 = 80.64516
## 21 = 82.7957
## 22 = 81.72043
## 23 = 82.7957
## 24 = 80.64516
## 25 = 81.72043
## 26 = 81.72043
## 27 = 81.72043
## 28 = 82.7957

```

```

#Accuracy plot
plot(k.optm, type="b", xlab="K- Value", ylab="Accuracy level")

```

